

Building applications for interactive data exploration in systems biology

Bjørn Fjukstad
Department of Computer Science,
UiT — The Arctic University of
Norway
Tromsø 9037
bjorn.fjukstad@uit.no

Vanessa Dumeaux
Department of Biology, Concordia
University
Montréal H4B 1R6
vanessadumeaux@gmail.com

Karina Standahl Olsen
Department of Community
Medicine, UiT — The Arctic
University of Norway
Tromsø 9037
karina.standahl.olsen@uit.no

Eiliv Lund
Department of Community
Medicine, UiT — The Arctic
University of Norway
Tromsø 9037
eiliv.lund@uit.no

Michael Hallett
Department of Biology, Concordia
University
Montréal H4B 1R6
michael.hallett@concordia.ca

Lars Ailo Bongo
Department of Computer Science,
UiT — The Arctic University of
Norway
Tromsø 9037
larsab@cs.uit.no

ABSTRACT

As the systems biology community generates and collects data at an unprecedented rate, there is a growing need for interactive data exploration tools to explore the datasets. These tools need to combine advanced statistical analyses, relevant knowledge from biological databases, and interactive visualizations in an application with clear user interfaces. To answer specific research questions tools must provide specialized user interfaces and visualizations. While these are application-specific, the underlying components of a data analysis tool can be shared and reused later. Application developers can therefore compose applications of reusable services rather than implementing a single monolithic application from the ground up.

Our approach for developing data exploration applications in systems biology builds on the microservice architecture. Microservice architectures separates an application into smaller components that communicate using language-agnostic protocols. We show that this design is suitable in bioinformatics applications where applications often use different tools written in different languages by different research groups. By deploying each service using software containers, it enables re-use and sharing of key components between applications, reducing development, deployment, and maintenance time.

We demonstrate the viability of our approach through a web application, MIXT blood-tumor, for exploring and comparing transcriptional profiles from blood and tumor samples in breast cancer patients. The application integrates

advanced statistical software, up-to-date information from biological databases, and modern data visualization libraries.

The web application for exploring transcriptional profiles, MIXT, is online at mixt-blood-tumor.bci.mcgill.ca and open-sourced at github.com/fjukstad/mixt. Packages to build the supporting microservices are open-sourced as a part of Kvik at github.com/fjukstad/kvik.

KEYWORDS

Interactive data exploration, containers, reproducible, microservice

ACM Reference format:

Bjørn Fjukstad, Vanessa Dumeaux, Karina Standahl Olsen, Eiliv Lund, Michael Hallett, and Lars Ailo Bongo. 2017. Building applications for interactive data exploration in systems biology. In *Proceedings of The 8th ACM Conference on Bioinformatics, Computational Biology, and Health Informatics (ACM BCB)*, Boston, Massachusetts USA, August 2017 (ACM-BCB'17), 9 pages. DOI: 10.475/123.4

INTRODUCTION

In recent years the biological community has generated an unprecedented amount of data. While the cost of data collection has drastically decreased, data analysis continues to be a larger fraction of the total experiment cost.[18] An important part of data analysis includes the time spent by human experts interpreting the results. This calls for novel methods for building data analysis tools for data exploration and interpretation.

In the field of systems biology, data exploration applications need to link results to relevant prior knowledge. In the later years there has been a tremendous effort to curate databases with relevant information on genes and processes. Databases such as the Molecular Signatures Database (MSigDB)¹ and the Kyoto Encyclopedia of Genes and

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

ACM-BCB'17, Boston, Massachusetts USA

© 2017 ACM. 978-1-4503-2894-4...\$15.00

DOI: 10.475/123.4

¹software.broadinstitute.org/gsea/msigdb

Genomes (KEGG)² both provide interfaces to retrieve data that can be used to better the understanding of data analysis results.

Several tools for biological data analysis are now available in various programming languages. These include a wide variety of bioinformatics methodologies and graphical analysis tools. In the R statistical programming language developers share software through repositories such as CRAN³ or Bioconductor⁴. In other languages, libraries for biological computation are often available like BioPython[2] and biogo[8] for Python or Go, respectively. Projects such as Galaxy⁵ and Common Workflow Language (CWL)⁶ enable researchers to build and run biological data analysis pipelines consisting of a wide range of tools. Although these frameworks are tremendously helpful, we need novel approaches to build applications that integrate high-performance bioinformatics tools with specialized user-interfaces and interactive visualizations

Different programming languages solve different tasks. For example, new biological data analysis techniques are quickly released in R and its package repositories; high-performance computer vision tasks are performed in C++ and OpenCV; and portable user interfaces more easily built in HTML, CSS and JavaScript. While it would make development easy, we argue that data exploration applications in systems biology should not be written in a single programming language. Therefore applications that integrate novel statistical analysis tools, interactive visualizations, and biological databases likely need to include several components written in different languages.

A microservice architecture structures an application into small reusable, loosely coupled parts. These communicate via lightweight programming language-agnostic protocols such as HTTP, making it possible to write single applications in multiple different programming languages. This way the most suitable programming language is used for each specific part. To build a microservice application, developers bundle each service in a container. Containers are built from configuration files which describe the operating system, software packages and versions of these. This makes reproducing the analyses, database lookups, library versions in an application a trivial task. The most popular implementation of a software container is Docker⁷, but others such as Rkt⁸ exist. Initiatives such as BioContainers⁹ now provide containers pre-installed with different bioinformatics tools.

From our experience we have identified a set of components and features we need to build data exploration applications.

- (1) A low-latency language-independent approach for integrating, or embedding, statistical software, such as R, directly in an application.
- (2) Low latency language-independent interface to online reference databases in biology that users can query to better understand results from statistical analyses.
- (3) A simple method for deploying and sharing the components of an application between projects.

In this paper, we describe a novel approach for building data exploration applications in systems biology. We show that by building applications as a set of services we can reuse and share its components between applications. The key services of a biological data exploration application are i) a compute service for executing statistical analyses in languages such as R, ii) a database query service for retrieving information from biological databases, and iii) the user-facing visualizations and user-interfaces. In addition, by packaging the services using container technology they are easy to deploy, simple to reproduce, and easy to share between projects. We have used our approach to build a number of applications, both command-line and web-based. In this paper we describe how we used our approach to develop MixT, a web application for exploring and comparing transcriptional profiles from blood and tumor samples.

RELATED WORK

In this section we discuss methods for integrating statistical analyses in data exploration applications, different visualization frameworks, interfaces to biological reference databases, and methods for containerizing applications.

Integrate Statistical Analyses

Shiny is a web application framework for R¹⁰. It allows developers to build web applications in R without having to have any knowledge about HTML, CSS or Javascript. Its widget library provides more advanced Javascript visualizations such as Leaflet¹¹ for maps or three.js¹² for WebGL-accelerated graphics. Developers can choose to host their own web server with the user-built Shiny Apps, or host them on public servers. Shiny forces users to implement data exploration applications in R, limiting the functionality to the widgets and libraries in Shiny.

OpenCPU is a system for embedded scientific computing and reproducible research.[14] Similar to the compute service in Kvik, it offers an HTTP API to the R programming language to provide an interface with statistical methods. It allows users to make function calls to any R package and retrieve the results in a wide variety of formats such as JSON or PDF. Users can choose to host their own R server or use public servers, and OpenCPU works in a single-user setting within an R session, or a multi-user setting facilitating multiple parallel requests. This makes OpenCPU suitable for building a service that can run statistical analyses. OpenCPU

²kegg.jp

³cran.r-project.org

⁴bioconductor.org

⁵galaxyproject.org

⁶commonwl.org

⁷docker.com

⁸coreos.com/rkt

⁹biocontainers.pro

¹⁰shiny.rstudio.com

¹¹leafletjs.com

¹²threejs.org

provides a Javascript library for interfacing with R, as well as Docker containers for easy installation and OpenCPU has been used to build multiple applications.¹³ To achieve better performance we chose to build our compute service rather than using OpenCPU. However, in Kvik we provide a Go package to interface with OpenCPU since it follows the same interface as our servers.

Renjin is a JVM-based interpreter for the R programming language.^[1] It enables integrating the R interpreter in web applications. Since it is built on top of the JVM it allows developers to write data exploration applications in Java that interact directly with R code, both running on top of the JVM. Although Renjin provides as a service pre-built versions of packages from CRAN and BioConductor, not all packages can be built for use in Renjin. This makes it necessary to re-implement the packages that cannot be built for use in Renjin.

Parallel and Distributed Execution

Biogo is a bioinformatics library written in Go. It provides functionality to analyze genomic and metagenomic datasets in the Go programming language.^[?] Using the Go programming language the developers are able to provide high-performance parallel processing in a clean and simple programming language. Go provides the ease of programming as popular scripting languages such as R, but the performance of low-level languages such as C++.

MapReduce is a popular framework and programming model to analyze large-scale datasets using distributed and parallel computing.^[3] While MapReduce has shown its usefulness in certain applications, it does not fit the different statistical methods in bioinformatics. An alternative to MapReduce is Spark, which supports high-performance parallel execution of iterative machine learning methods by partitioning data across a set of machines and optimizing the data access.^[20] Spark is growing in popularity, especially with the Spakrlyr¹⁴ and SparkR¹⁵ R packages that allow R users to run analyses on top of Spark without having to port the analysis code to another language.

ADAM and VariantSpark are systems that are implemented on top of Spark to analyse genomic data. ADAM provides a set of formats, APIs and processing stage implementations^[12] while VariantSpark provides an method for clustering genomic variants^[15].

Pachyderm¹⁶ is a system for running containerized data analysis pipelines. Each step in an analysis pipeline is run within a software container and the output data is version controlled. Pachyderm automatically partitions and distributes the input data to enable parallel processing. Pachyderm runs on top of Kubernetes¹⁷, a system for managing and deploying containerized applications.

Visualization tools

Cytoscape is an open source software platform for visualizing complex networks and integrating these with any type of attribute data^[19]. It allows for analysis and visualization in the same platform. Users can add additional features, such as databases connections or new layouts, through Apps. One such app is cyREST which allows external network creation and analysis through a REST API^[13]. To bring the visualization and analysis capabilities to the web the creators of Cytoscape have developed Cytoscape.js¹⁸, a Javascript library to create interactive graph visualizations.

Caleydo is a framework for building applications for visualizing and exploring biomolecular data^[?]. Until 2014 it was a standalone tool that needed to be downloaded, but the Caleydo team are now making the tools web-based. There have been several applications built using Caleydo: StratomeX for exploring stratified heterogeneous datasets for disease subtype analysis^[11]; Pathfinder for exploring paths in large multivariate graphs^[16]; UpSet to visualize and analyse sets, their intersections and aggregates^[9]; Entourage and enRoute to explore and visualize biological pathways^{[10][17]}; LineUp to explore rankings of items based on a set of attributes^[7]; and Domino for exploring subsets across multiple tabular datasets^[6].

BioJS is an open-source JavaScript framework for biological data visualization.^[5] It provides a community-driven online repository with a wide range components for visualizing biological data contributed by the bioinformatics community. BioJS builds on node.js¹⁹ providing both server-side and client-side libraries.

Biological Databases

Maybe some words here on how to get data out of the different biological databases? + possible provenance.

Containerized analysis

In the later years software containers have been widely adopted by both the software industry as well as research communities. Containers provide an isolated execution environment that can be used to package and run an application with all its dependencies, library versions and configuration files. In researchers containers have become popular because they provides a reproducible environment that can be shared between research projects.

In bioinformatics researchers are starting to bundle their software using software containers, such as Docker. There are repositories such as BioContainers^[?] and BioBoxes^[?] that provide containers preinstalled with software for doing analyses and running different applications. Systems such as Galaxy now allows researchers to build analysis pipelines where each step is executed within a software container.

¹³ opencpu.org/apps.html

¹⁴ spark.rstudio.com

¹⁵ spark.apache.org/docs/latest/sparkr.html

¹⁶ pachyderm.io

¹⁷ kubernetes.io

¹⁸ js.cytoscapejs.org

¹⁹ nodejs.org

Kvik and Kvik Pathways

We have previously built a system for interactively exploring gene expression data in context of biological pathways.[4] Kvik Pathways is a web application that integrates gene expression data from the Norwegian Women and Cancer (NOWAC) cohort together with pathway images from the Kyoto Encyclopedia of Genes and Genomes (KEGG). We used the experience building Kvik Pathways to completely re-design and re-implement the R interface in Kvik. From having an R server that can run a set of functions from an R script, it now has a clean interface to call any function from any R package, not just retrieving data as a text string but in a wide range of formats. We also re-built the database interface, which is now a separate service. This makes it possible to leverage its caching capabilities to improve latency. This transformed the application from being a single monolithic application into a system that consists of a web application for visualizing biological pathways, a database service to retrieve pathway images and other metadata, and a compute service for interfacing with the gene expression data in the NOWAC cohort. We could then re-use the database and the compute service in the MlXt application.

METHODS

In this section we first motivate our microservice approach based on our experiences developing the MlXt web application. We describe the process from initial data analysis to the final application, highlighting the importance of language-agnostic services to facilitate the use of different tools in different parts of the application. We show that these services are easy to deploy and provide performance to use them to build data exploration applications. We then generalize the ideas to a set of principles and services that can be reused and shared between applications, and show their design and implementation.

Motivating Example

The aim of the Matched Interactions Across Tissues (MlXt) study was to identify genes and pathways in the primary breast tumor that are tightly linked to genes and pathways in the patient blood cells. We generated and analyzed expression profiles from blood and matched tumor cells in 173 breast cancer patients included in the Norwegian Women and Cancer (NOWAC) study.[?] The MlXt analysis starts by identifying sets of genes tightly co-expressed across all patients in each tissue. Each group of genes or modules were annotated based on known a priori biological knowledge about gene functionality. Focus was placed on the relationships between tissues by asking if specific biologies in one tissue are linked with (possibly distinct) biologies in the second tissue, and this within different subgroup of patients (i.e. subtypes of breast cancer).

We built an R package²⁰ with the statistical methods and static visualizations for identifying associations between modules across tissues. The exploration of the results encompass the examination of ~ 20 modules and their functional enrichments. That is $\sim 23 \times 19 = 437$ associations computed for each of the 22 patient subgroups.

To explore this large amount of data and results we needed an interactive point-and-click application that could interface with the statistical methods and link the results to online databases. This would make it possible for users to explore the results without any coding background. We needed an application that could interface with MSigDB to fetch gene set meta-data and the Entrez Programming Utilities (Eutils) to retrieve gene meta-data. To make the interfaces to the statistical analyses and databases re-usable by other applications later it was necessary for these to communicate using open protocols.

In addition, we required that the application was containerized. This allows us to deploy the application on a wide range of hardware, from local installations to deployments to cloud providers such as Amazon Web Services (AWS)²¹.

Design

Our experience can be generalized into the following design principles for building applications in bioinformatics:

Principle 1: Build applications as collections of language-agnostic microservices. This enables re-use of components and does not enforce any specific programming language on the user-facing logic or the underlying components of the application. Within bioinformatics, researchers and application developers use a wide range of programming languages to build tools. By composing an application of services communicating over standard protocols it is possible to re-use existing tools in new applications.

Principle 2: Use software containers to package each service. This has a number of benefits; it simplifies deployment, managing application dependencies and libraries is no longer a problem, and it simplifies sharing of services between developers.

Using these design principles we built the MlXt web application using the microservices in Kvik to interface with statistical analyses and biological databases. We built a compute service for executing statistical analyses and a database service for retrieving relevant information on genes and biological processes. Using these it is possible to develop specialized data exploration application in any modern programming language. In the rest of the section we discuss the design choices we made to build the microservices that power the MlXt web application.

Compute Service. The main cornerstone of every data exploration application in systems biology is a dataset. Datasets go through a series of transformations before they can be interpreted by experts. Because of the complexity of these datasets, package repositories such as BioConductor provide

²⁰github.com/vdumeaux/mixtR

²¹aws.amazon.com

software for reading and analyzing the data. Because of the complex datasets there are challenges in building systems for exploring them: i) datasets require specialized software to be read, ii) datasets can be too big to fit on a desktop computer, iii) statistical methods for analyzing the datasets are too computationally intensive to run on a desktop computer, iv) users may want to modify statistical analyses directly from an application, and v) users need to know exactly what transformations has been done to a dataset to reproduce the analyses.

From these challenges we have designed the compute service in Kvik. The compute service interfaces directly with the R programming language, making it possible to call functions from any R package. Application developers can use the compute service to execute analyses and return results. By interfacing directly with R developers can leverage this to produce dynamic applications. For example, if an application uses a clustering method to color nodes in a graph, end-users can tweak parameters that interactively within the application that changes the node coloring in graph visualization. Also, by interfacing with R directly from an application it can store provenance data on the statistical methods and input parameters being used. By placing data storage and analysis into a service, application developers can deploy it on a powerful server while keeping the user-facing application logic on the desktop. Packaging the service into a software container also provides the necessary functionality to ensure that the statistical analyses can be reproduced later.

The compute service offers three main operations to interface with R. i) to call a function from an R package, ii) to get the results from a previous function call, and iii) a catch-all term that both calls a function and returns the results. We use the same terminology as OpenCPU and have named the three operations Call, Get, and RPC respectively. These three operations provide the necessary interface for applications to include data in the applications.

Database Service. To understand analysis results experts query databases and scientific literature. There are a wealth of online databases, some of which provide open APIs in addition to web user interfaces that application developers can make use of. While the databases can provide helpful information, there are some challenges including them in interactive data exploration applications:

i) the APIs aren't fast enough to use in interactive applications where the application has to perform multiple database calls, and ii) some databases put restrictions on the number of database calls.

To solve these problems we built a database service for application developers to include. The service includes a simple caching mechanism that solves the two challenges. If we cache queries to the database we can speed up subsequent calls and reduce the load on the respective databases. Both the query from the application and the response from the databases are stored for later use. The database service provides an interface to biological databases for retrieving meta-data on

genes and processes. We have currently packages for interfacing with E-utilities²², MSigDB²³, Hugo Gene Nomenclature Committee (HGNC)²⁴, and Kyoto Encyclopedia of Genes and Genomes (KEGG)²⁵.

Applications. Kvik provides services to perform database lookup and execute statistical analyses. Since we provide both services as software containers, application developers simply run these and write applications that interface with them. Using our services application developers have a platform for running statistical analyses, and a fast clean interface to get metadata about genes and processes.

Figure 1 shows how the the Mlxt application is built using Kvik microservices. In Mlxt we built a specialized web application that interfaces with Kvik to get data from biological databases and to run statistical analyses from the mixt R package²⁶.




Figure 1: An overview of the relationship between the Mlxt application and Kvik. Mlxt contains a web application (online at mixt-blood-tumor.bci.mcgill.ca) and the R Package that provides analyses and data to the web application. Kvik provides the services for running the statistical analyses from the R package, and the database lookups found in the web application.

Implementation

In this section we describe the implementation details in Kvik its microservices.

Kvik is implemented as a collection of Go packages required to build services that can integrate statistical software in a data exploration and provide an interface to up-to-date biological databases. We chose the Go programming language

²²eutils.ncbi.nlm.nih.gov

²³software.broadinstitute.org/gsea/msigdb

²⁴genenames.org

²⁵kegg.jp

²⁶github.com/vdumeaux/mixtR

because of its performance, ease of development, and simple deployment. To integrate R we provide two packages *gopencpu* and *r*, that interface with OpenCPU and Kvik R servers respectively. To interface with biological databases we provide the packages *eutils*, *gsea*, *genenames*, and *kegg* that interface with E-utils, MsigDB, HGNC and KEGG respectively. In addition to these packages we provide Docker images that implement the two required microservices.

Both the compute and the databases service in Kvik builds on the standard *http* package in Go. The database service use the *gocache*²⁷ package to cache any query to an online database. In addition we deploy each service as Docker containers.²⁸

Compute Service. The compute service is an HTTP server that communicates with a pre-set number of R processes to execute statistical analyses. On start of the compute service launches a user-defined number of R worker sessions for executing analyses, default is 5. The compute service uses a round-robin scheduling scheme to distribute incoming requests to the workers. We provide a simple FIFO queue for queuing of requests. The compute service also provides the opportunity for users to cache analysis results to speed up subsequent calls.

The compute service in Kvik is built using a hybrid state pattern. A hybrid state pattern origins from functional programming, where output from a method only depends on its inputs and not the program state.[14]. In practice this means that the compute services stores the output results from function calls, but can compute them if the service has to restart. This makes it possible for us to scale the compute service horizontally to handle more requests. If an R worker session for some reason crashes the compute service simply starts up a replacement.

RESULTS

We show the viability of the microservices approach in Kvik by describing the MlXT web application for exploring and comparing transcriptional profiles from blood and tumor samples.

MlXT Analysis Tasks

We define six data exploration tasks that the web application should help users perform:

Explore co-expression gene sets in tumor and blood tissue. We want to simplify the process of exploring the computed co-expression gene sets, or modules, through the web-application. The application should visualize gene expression patterns together with clinicopathological variables for each module. In addition we want to enable users to study the underlying biological functions of each module by including gene set analyses between the module genes and known gene sets.

Explore co-expression relationships between genes. Users should be able to explore the co-expression relationship as a graph visualization. The network should visualize each gene as a node and a significant co-expression relationship as an edge.

Explore relationships between modules from each tissue. Users should be able to explore the relationship between modules from different tissues. We provide two different metrics to compare modules, and the web application should enable users to interactively browse these relationships. In addition to providing visualizations the compare modules from each tissue, users should also be able to explore the relationships, but for different breast cancer patient groups.

Explore relationships between clinical variables and modules. In addition to comparing the association between modules from both tissues, users also have the possibility of exploring the association with a module and a specific clinical variable. This should also be possible for the different breast cancer patient group.

Explore association between user-submitted gene lists and computed modules. We want to enable users to explore their own gene lists to explore them in context of the co-expression gene sets. The web application must handle uploads of gene lists and compute association between the genelist and the MlXT modules on demand.

Search for genes or gene lists of interest. To facilitate faster lookup of genes and biological processes, the web application should provide a search functionality that lets users locate genes or gene lists and show association to the co-expression gene sets.

MlXT Design and Implementation

From these six analysis tasks we designed and implemented MlXT as a web application that integrates statistical analyses and information from biological databases together with interactive visualizations. The MlXT web application consists of three services: i) the web application itself containing the user-interface and visualizations; ii) the compute service performing the MlXT analyses delivering data to the web application; and iii) the database service providing up-to-date information from biological databases. Each of these services run within Docker containers making the process of deploying the application simple.

We structured the MlXT application with a separate view for each analysis task. To explore the co-expression gene sets we built a view that combines both static visualizations from R together with interactive tables with gene overlap analyses. Figure 2 shows the web page presented to users when they access the co-expression gene set 'darkturquoise' from blood. Using the Kvik compute service we can generate plots on demand and provide users with high-resolution PDFs or PNG files.

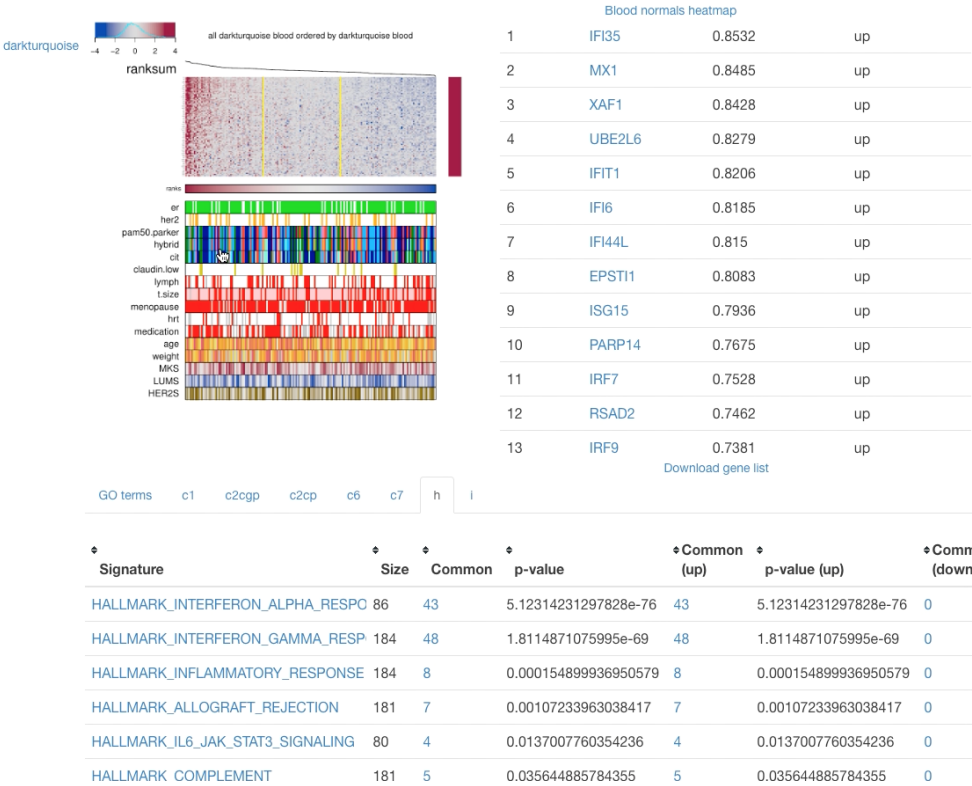
To explore the co-expression relationship between genes we use an interactive graph visualization build with Sigma.js²⁹. We have built visualization for both tissues, with graph sizes

²⁷ github.com/fjukstad/gocache

²⁸ Available at hub.docker.com/r/fjukstad/kvik-r and hub.docker.com/r/fjukstad/db

²⁹ sigmajs.org

Figure 2: MlXt module overview page. The screenshot show the user interface for exploring a single module. It consists of three panels. The top left panel contains the gene expression heatmap. The top right panel contains a table of the genes found in the module. The bottom panel contains the results of gene overlap analyses from the module genes and known gene sets from MSigDB.




of 2705 nodes and 90 348 edges for the blood network, and 2066 nodes and 50 563 edges for the biopsy network. The sigma.js visualization library has functionality for generating a layout for large networks, but we generate this layout server-side to reduce the computational load on the client. To generate this layout we use the GGally package³⁰. By generating the network layout using the compute service we relieve the clients.

To visualize relationships between modules from different tissues, or their relationship to clinical variables we built a heatmap visualization using the d3³¹ library. Figure 3 shows an example of this heatmap visualization, showing the association between the clinical variables and the modules from biopsy for all samples.

Evaluation

We evaluate the performance of the services by running a set of benchmarks. To evaluate the database service we measure the query time for retrieving information about a specific gene with and without caching. This illustrates how we can improve performance in an application by using a database

Figure 3: Heatmap visualization of the association between clinical variables and the modules in biosy. The visualization is built using the d3 JavaScript library.



service rather than accessing the database directly. We also measure storage requirements for the caching mechanism to show that it does not require a large storage capacity. To

³⁰cran.r-project.org/web/packages/GGally

³¹d3.js.org

show that it scales we show the query time for 1, 10, 100, and 1000 concurrent requests.

We evaluate the compute service by performing two operations, first generate a set of numbers then plot them and return the resulting visualization. We show that the latency is low enough to use it in an interactive application. To show that it performs under heavy load we perform the same operations using 1, 10, 100, and 1000 concurrent requests.

TODO: INSERT NUMBERS

DISCUSSION

The compute service in Kvik follows many of the design patterns in OpenCPU. Both systems interface with R packages using a hybrid state pattern over HTTP. Both systems provide the same interface to execute analyses and retrieve results. While OpenCPU is implemented on top of R and Apache, Kvik is implemented from the ground up in Go. Because of the similarities in the interface to R in Kvik we provide packages for interfacing with our own R server or OpenCPU R servers through the *gopencpu* package.³²

Future work

While we have built

We have built a database service that provides a sufficient interface for the MixT web application. While we have developed the software packages for interfacing with more databases, these haven't been included in the database service yet.

One large concern that we haven't addressed in this paper is security. In particular one security concern that we plan to address in Kvik is the restrictions on

Docker compose

This includes both access restrictions to parts of an application, e.g. access to a database, and security concerns regarding remote code execution within the compute service. to the biological databases that are

Security. Documentation. Scale.

CONCLUSIONS

We have designed an approach for building data exploration applications in systems biology that builds on a microservice architecture. Using this approach we have built a web application that leverages this architecture to integrate statistical analyses, interactive visualizations and data from biological databases.

LIST OF ABBREVIATIONS USED

COMPETING INTERESTS

The authors declare that they have no competing interests.

REFERENCES

- [1] Alexander Bertram. 2013. Renjin: The new R interpreter built on the JVM. In *The R User Conference, user! 2013 July 10-12 2013 University of Castilla-La Mancha, Albacete, Spain*, Vol. 10. 105.
- [2] Peter JA Cock, Tiago Antao, Jeffrey T Chang, Brad A Chapman, Cymon J Cox, Andrew Dalke, Iddo Friedberg, Thomas Hamelryck, Frank Kauff, Bartek Wilczynski, and others. 2009. Biopython: freely available Python tools for computational molecular biology and bioinformatics. *Bioinformatics* 25, 11 (2009), 1422–1423.
- [3] Jeffrey Dean and Sanjay Ghemawat. 2008. MapReduce: simplified data processing on large clusters. *Commun. ACM* 51, 1 (2008), 107–113.
- [4] Bjørn Fjukstad, Karina Standahl Olsen, Mie Jareid, Eiliv Lund, and Lars Ailo Bongo. 2015. Kvik: three-tier data exploration tools for flexible analysis of genomic data in epidemiological studies. *F1000Research* 4 (2015).
- [5] John Gómez, Leyla J García, Gustavo A Salazar, Jose Villaveces, Swanand Gore, Alexander García, Maria J Martín, Guillaume Launay, Rafael Alcántara, Noemi Del Toro Ayllón, and others. 2013. BioJS: an open source JavaScript framework for biological data visualization. *Bioinformatics* (2013), btt100.
- [6] Samuel Gratzl, Nils Gehlenborg, Alexander Lex, Hanspeter Pfister, and Marc Streit. 2014. Domino: Extracting, comparing, and manipulating subsets across multiple tabular datasets. *IEEE transactions on visualization and computer graphics* 20, 12 (2014), 2023–2032.
- [7] Samuel Gratzl, Alexander Lex, Nils Gehlenborg, Hanspeter Pfister, and Marc Streit. 2013. Lineup: Visual analysis of multi-attribute rankings. *IEEE transactions on visualization and computer graphics* 19, 12 (2013), 2277–2286.
- [8] R Daniel Kortschak and David L Adelson. 2014. bfo: a simple high-performance bioinformatics toolkit for the Go language. *bioRxiv* (2014). DOI: <https://doi.org/10.1101/005033> arXiv: <http://biorxiv.org/content/early/2014/05/12/005033.full.pdf>
- [9] Alexander Lex, Nils Gehlenborg, Hendrik Strobelt, Romain Vuillemot, and Hanspeter Pfister. 2014. UpSet: visualization of intersecting sets. *IEEE transactions on visualization and computer graphics* 20, 12 (2014), 1983–1992.
- [10] Alexander Lex, Christian Partl, Denis Kalkofen, Marc Streit, Samuel Gratzl, Anne Mai Wassermann, Dieter Schmalstieg, and Hanspeter Pfister. 2013. Entourage: Visualizing relationships between biological pathways using contextual subsets. *IEEE transactions on visualization and computer graphics* 19, 12 (2013), 2536–2545.
- [11] Alexander Lex, Marc Streit, H-J Schulz, Christian Partl, Dieter Schmalstieg, Peter J Park, and Nils Gehlenborg. 2012. StratomeX: Visual Analysis of Large-Scale Heterogeneous Genomics Data for Cancer Subtype Characterization. In *Computer graphics forum*, Vol. 31. Wiley Online Library, 1175–1184.
- [12] Matt Massie, Frank Nothaft, Christopher Hartl, Christos Kozaniotis, André Schumacher, Anthony D Joseph, and David A Patterson. 2013. Adam: Genomics formats and processing patterns for cloud scale computing. *University of California, Berkeley Technical Report, No. UCB/EECS-2013 207* (2013).
- [13] Keiichiro Ono, Tanja Muetze, Georgi Kolishovski, Paul Shannon, and Barry Demchak. 2015. CyREST: Turbocharging Cytoscape Access for External Tools via a RESTful API. *F1000Research* 4 (2015).
- [14] Jeroen Ooms. 2014. The OpenCPU System: Towards a Universal Interface for Scientific Computing through Separation of Concerns. *arXiv preprint arXiv:1406.4806* (2014).
- [15] Aidan R O'Brien, Neil FW Saunders, Yi Guo, Fabian A Buske, Rodney J Scott, and Denis C Bauer. 2015. VariantSpark: population scale clustering of genotype information. *BMC genomics* 16, 1 (2015), 1052.
- [16] Christian Partl, Samuel Gratzl, Marc Streit, Anne Mai Wassermann, Hanspeter Pfister, Dieter Schmalstieg, and Alexander Lex. 2016. Pathfinder: Visual Analysis of Paths in Graphs. In *Computer Graphics Forum*, Vol. 35. Wiley Online Library, 71–80.
- [17] Christian Partl, Alexander Lex, Marc Streit, Denis Kalkofen, Karl Kashofer, and Dieter Schmalstieg. 2012. enRoute: Dynamic path extraction from biological pathway maps for in-depth experimental data analysis. In *Biological Data Visualization (BioVis), 2012 IEEE Symposium on*. IEEE, 107–114.
- [18] Andrea Sboner, Ximeng Jasmine Mu, Dov Greenbaum, Raymond K Auerbach, and Mark B Gerstein. 2011. The real cost of sequencing: higher than you think! *Genome biology* 12, 8 (2011), 125.
- [19] Paul Shannon, Andrew Markiel, Owen Ozier, Nitin S Baliga, Jonathan T Wang, Daniel Ramage, Nada Amin, Benno

³² github.com/fjukstad/kvik/tree/master/gopencpu

- Schwikowski, and Trey Ideker. 2003. Cytoscape: a software environment for integrated models of biomolecular interaction networks. *Genome research* 13, 11 (2003), 2498–2504.
- [20] Matei Zaharia, Mosharaf Chowdhury, Michael J Franklin, Scott Shenker, and Ion Stoica. 2010. Spark: Cluster Computing with Working Sets. *HotCloud* 10, 10-10 (2010), 95.